

End-to-end latency: py3.12 + adaptive cuML vs py3.8 baseline

Datacenter GPUs — orange = py3.12 slower, green = py3.12 faster (values above axis capped; real value annotated)

