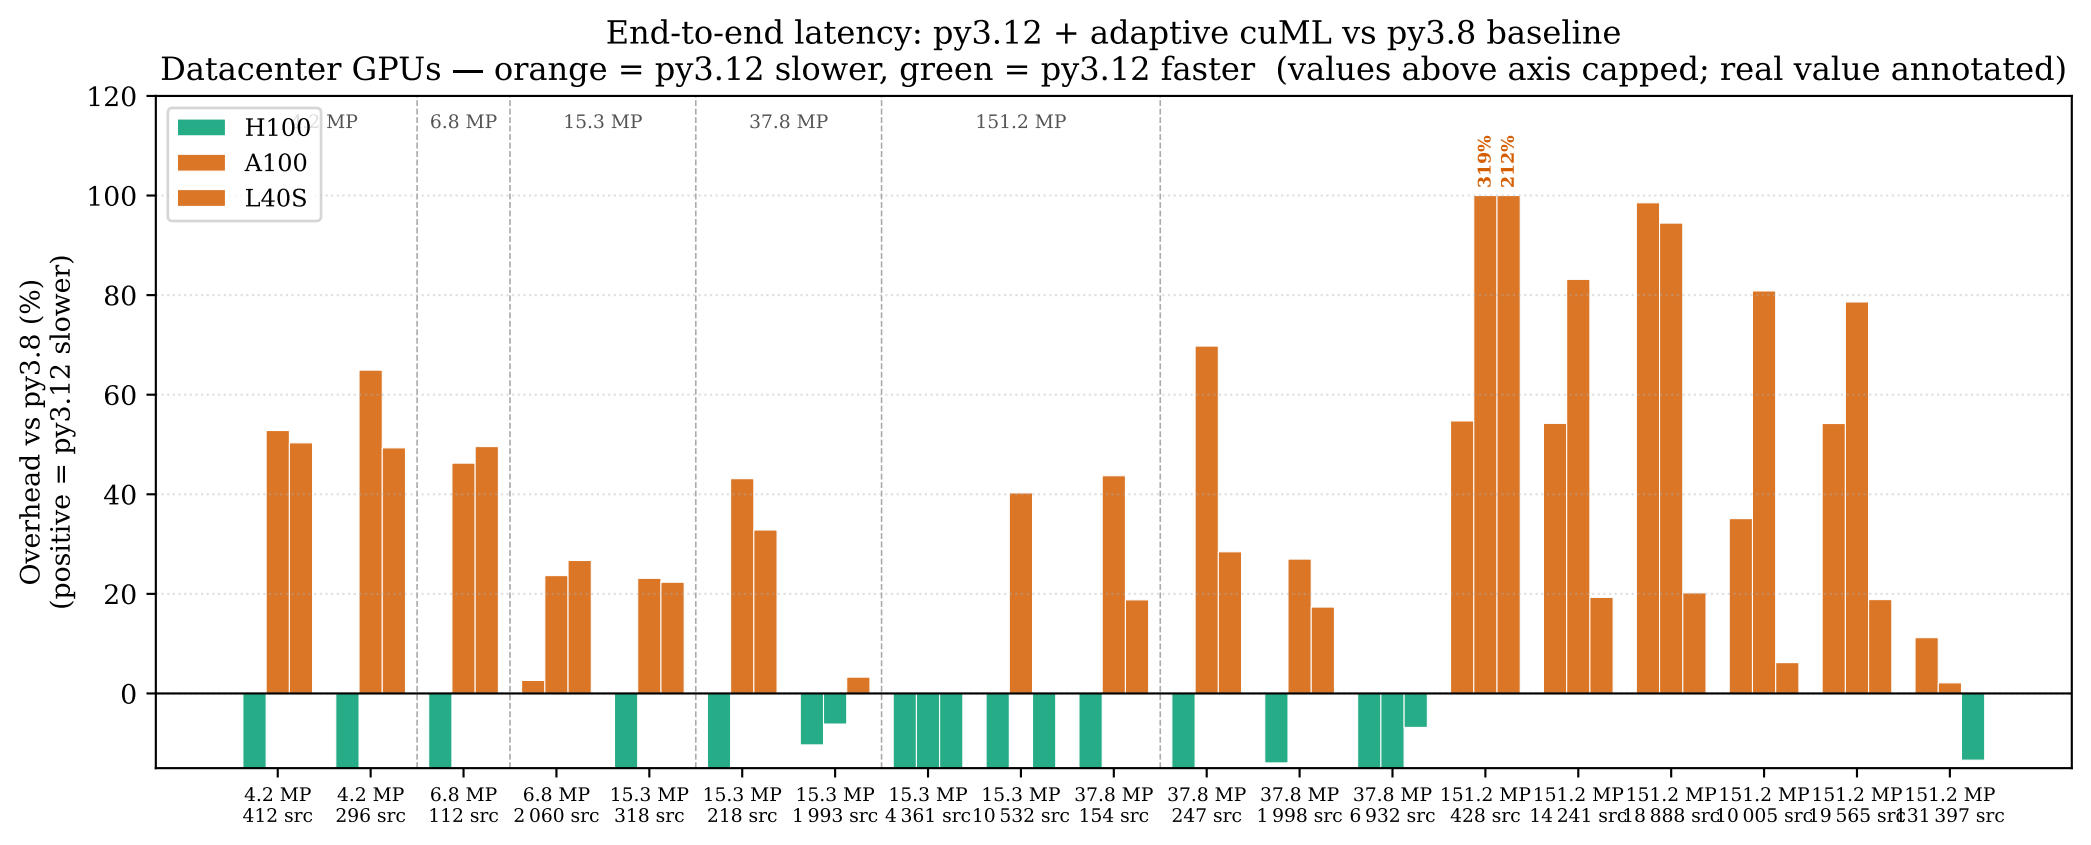

End-to-end latency: py3.12 + adaptive cuML vs py3.8 baseline
Datacenter GPUs — orange = py3.12 slower, green = py3.12 faster (values above axis capped; real value annotated)
Overhead vs py3.8 (%)
(positive = py3.12 slower)
H100
A100
L40S
6.8 MP
15.3 MP
37.8 MP
151.2 MP
319%
212%
120
100
80
60
40
20
0
4.2 MP 412 src
4.2 MP 296 src
6.8 MP 112 src
6.8 MP 2 060 src
15.3 MP 318 src
15.3 MP 218 src
15.3 MP 1 993 src
15.3 MP 4 361 src
15.3 MP 10 532 src
37.8 MP 154 src
37.8 MP 247 src
37.8 MP 1 998 src
37.8 MP 6 932 src
151.2 MP 428 src
151.2 MP 14 241 src
151.2 MP 8 888 src
151.2 MP 10 005 src
151.2 MP 19 565 src
151.2 MP 31 397 src