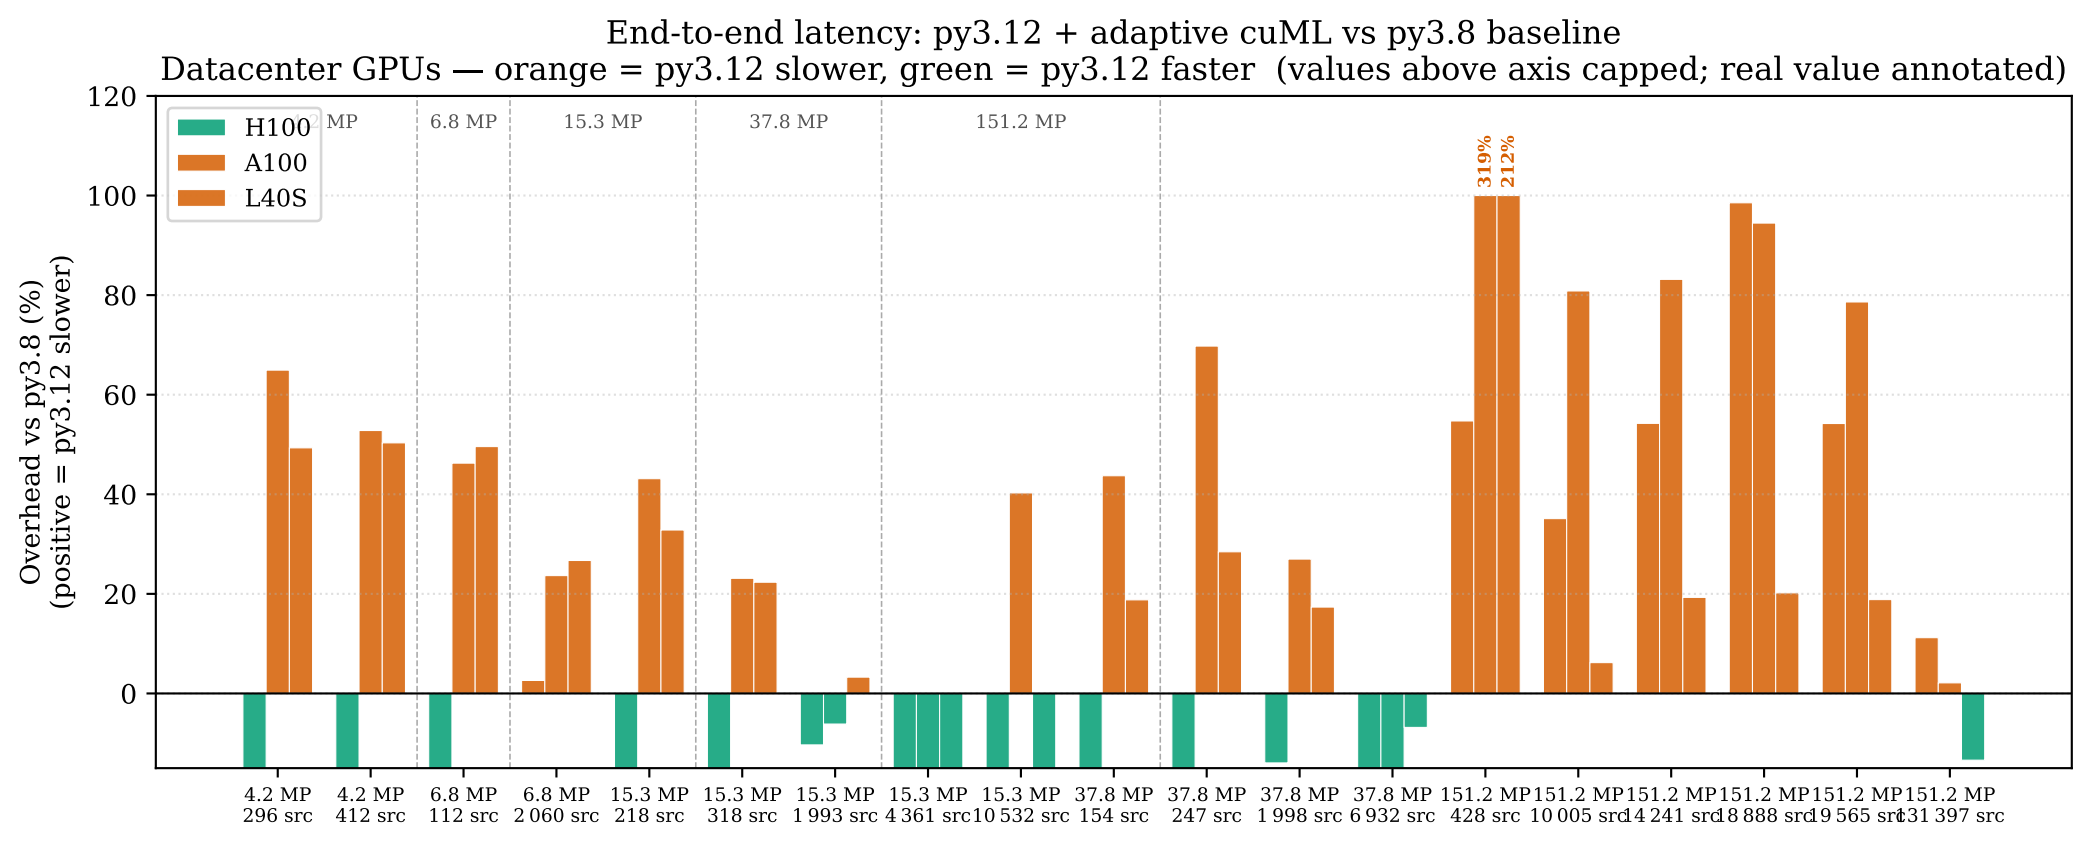

End-to-end latency: py3.12 + adaptive cuML vs py3.8 baseline
Datacenter GPUs — orange = py3.12 slower, green = py3.12 faster (values above axis capped; real value annotated)
Overhead vs py3.8 (%)
(positive = py3.12 slower)
H100
A100
L40S
6.8 MP
15.3 MP
37.8 MP
151.2 MP
319%
212%
4.2 MP 296 src
4.2 MP 412 src
6.8 MP 112 src
6.8 MP 2 060 src
15.3 MP 218 src
15.3 MP 318 src
15.3 MP 1 993 src
15.3 MP 4 361 src
15.3 MP 10 532 src
37.8 MP 154 src
37.8 MP 247 src
37.8 MP 1 998 src
37.8 MP 6 932 src
151.2 MP 428 src
151.2 MP 10 005 src
151.2 MP 14 241 src
151.2 MP 18 888 src
151.2 MP 19 565 src
151.2 MP 31 397 src